Nishal Stanislaus Silva, *PhD*

AI Forward Deployed Engineer | Production ML Deployment, Real-Time Inference, Customer-Facing Delivery

Canadian Permanent Resident, Eligible to work in Canada without sponsorship

nishal.silva@hotmail.com | +1 613 200 2030 | Toronto, ON | nishal.xyz | linkedin.com/in/nishal-silva

---

14th October 2026

Re: AI Inference Engineer, Forward Deployed Engineering – Baseten

Dear Hiring Manager,

I am applying for the Forward Deployed Engineer role at Baseten. I am a machine learning engineer with a PhD (University of Trento, 2025) and 10+ years across industry and academia, and most of that work has been taking an unclear problem to a production service that someone depends on. At MAS Holdings I led end-to-end development and deployment of computer vision systems in live manufacturing sites, improving operational efficiency by up to 300%.

The role asks for ownership from problem framing to evaluation, deployment and monitoring. That is how I worked in my postdoc: I owned the full ML pipeline for streaming sensor data from acquisition to deployment with monitoring, and built real-time inference services under 30 ms. The published version of this work runs at 14 ms on a Raspberry Pi 4, 74x faster than the baseline (F1 0.76). At MAS and at Forestpin I worked directly with operations, compliance and risk stakeholders to turn vague needs into specs, milestones and measurable results. My main language is Python, and I have deployed with Docker, AWS and CI/CD.

I do not have a formal pre-sales or customer-success title. My customer-facing experience comes from consulting and stakeholder-driven delivery, and I would be learning the sales and expansion side of the job. I am comfortable writing specs, running proofs of concept and owning a project when the requirements are still moving.

Baseten runs inference for companies that cannot afford it to be slow or unreliable, and that is the problem I have spent years on under tight latency and compute limits. I would like to help your customers get their models into production, measured on quality, latency and cost.

I am a Canadian permanent resident based in Toronto, ON, eligible to work without sponsorship, and open to hybrid or remote work from Toronto, or relocation within Canada. I am available to start with short notice. I would welcome the opportunity to discuss how my experience in production ML deployment and real-time systems can support the Forward Deployed Engineering team at Baseten. My portfolio and publications are at nishal.xyz.

Sincerely,

Nishal Stanislaus Silva, PhD