# Nishal Stanislaus Silva, *PhD*

Machine Learning Engineer | Computer Vision, Real-Time Inference, Cloud & Edge Deployment | Applied ML

*Canadian Permanent Resident, Eligible to work in Canada without sponsorship*

nishal.silva@hotmail.com | +1 613 200 2030 | Toronto, ON | nishal.xyz | linkedin.com/in/nishal-silva


---

18th October 2026

**Re: AI Computer Vision Engineer – AssistIQ**

Dear Hiring Manager,

I am applying for the AI Computer Vision Engineer role at AssistIQ. I am a Machine Learning Engineer with a PhD (University of Trento, 2025) and 10+ years across industry and academia, and I build end-to-end ML systems that run in production. At MAS Holdings I led computer vision systems for manufacturing, including care-label inspection that cut inspection time by 99.5% and raised operational efficiency by up to 300%. I also built a computer vision device for COVID-19 patient vitals that was deployed across hospital wards in Sri Lanka and recognised by the Government Medical Officers' Association. Using vision to give healthcare teams accurate data is work I want to keep doing, and AssistIQ's supply tracking mission is a good place to do it.

Three parts of your posting match my work directly. For classical computer vision, I built OpenCV-based inspection systems at MAS, taking a care-label checker from a flatbed scanner to an industrial camera with a conveyor and PLC reject, and a loom-side fabric defect detector on a camera array with a few-second alert budget. For real-time inference and edge deployment, my published system runs at 14 ms on a Raspberry Pi 4, 74x faster than the DTW baseline (Journal of the Audio Engineering Society, 2026), and I owned its full lifecycle from data acquisition to deployment and monitoring. For evaluation, I design benchmark suites with baselines and ablations, and I released four open datasets on Zenodo. I also use Claude Code as a core part of my development workflow.

I should be direct about two gaps. My production vision work has been classical, OpenCV-based pipelines, and my deep learning work has been mostly on sequence models and multimodal sensor data, so I have not yet shipped CNN-based detectors or vision-language models. I have worked with PyTorch and HuggingFace Transformers and would build your VLM evaluation on the same benchmarking approach. My MLOps experience is Airflow, Docker, MLflow and CI/CD on AWS, and I have not run a large-scale monitoring stack for a vision product. I have built the same pieces with other tools and expect to pick up yours quickly.

I am a Canadian permanent resident based in Toronto, ON, eligible to work without sponsorship, and available immediately. Remote work suits me well. I would welcome the opportunity to discuss how my experience in computer vision and real-time ML systems can support the AssistIQ computer vision team. My portfolio, publications and open datasets are available at nishal.xyz.

Sincerely,

**Nishal Stanislaus Silva**, PhD