Nishal Stanislaus Silva, *PhD*

Machine Learning Research Engineer | Multilingual NLP, Large-Scale Data Pipelines, Deep Learning | Applied ML

Canadian Permanent Resident, Eligible to work in Canada without sponsorship

nishal.silva@hotmail.com | +1 613 200 2030 | Toronto, ON | nishal.xyz | linkedin.com/in/nishal-silva

---

18th October 2026

Re: Member of Technical Staff, Multilingual – Cohere

Dear Hiring Manager,

I am applying for the Member of Technical Staff, Multilingual role at Cohere. I am a Machine Learning Research Engineer with a PhD (University of Trento, 2025) and 10+ years across industry and academia, and I build end-to-end ML pipelines from data acquisition through evaluation and deployment. Cohere's Multilingual team improving language models for people regardless of language is work I want to contribute to, and the role asks for large-scale data processing, strong Python and rigorous evaluation, which is how I have worked throughout my research.

Three parts of my work are relevant. For ML pipelines and engineering, I owned full pipelines for streaming multi-dimensional data in the EU-funded MUSMET project, and my published system runs at 14 ms on a Raspberry Pi 4, 74x faster than the DTW baseline, at F1 0.76 (Journal of the Audio Engineering Society, 2026). For evaluation and data, I design benchmark suites with baselines and ablations, and I released four open datasets on Zenodo with 9,800+ recordings. For multilingual work, I built speakfrench, which benchmarked six pronunciation-scoring methods, including Wav2Vec2 XLSR-FR, on 10,000 sentences and 100,000 scored trials (best ROC-AUC 0.822 against 0.599 for a DTW baseline), and I started poetry2music, a multilingual English, Italian, French and Sinhala annotated poetry corpus.

I should be direct about the gaps. My NLP and LLM experience is applied rather than research-level, and I have not trained large language models at scale. My multilingual work is on speech and poetry corpora, not LLM training data. I speak English natively and Sinhala as my mother tongue, with Italian at B1 and French at about A2. I have 10+ peer-reviewed publications in audio and signal processing venues, and I expect to ramp up quickly on multilingual LLM data and evaluation.

I am a Canadian permanent resident based in Toronto, ON, eligible to work without sponsorship, and available immediately. I would welcome the opportunity to discuss how my experience can support the Multilingual team at Cohere. My portfolio, publications and open datasets are available at nishal.xyz.

Sincerely,

Nishal Stanislaus Silva, PhD