Nishal Stanislaus Silva, *PhD*

Machine Learning Engineer | End-to-End ML Systems, Model Serving, MLOps | Applied ML

Canadian Permanent Resident, Eligible to work in Canada without sponsorship

nishal.silva@hotmail.com | +1 613 200 2030 | Toronto, ON | nishal.xyz | linkedin.com/in/nishal-silva

---

19th October 2026

Re: AI/ML Engineer – lululemon

Dear Hiring Manager,

I am applying for the AI/ML Engineer role on the Enterprise Data & AI team at lululemon. I am a Machine Learning Engineer with a PhD (University of Trento, 2025) and 10+ years across industry and academia, and I build ML systems end to end: dataset curation, feature engineering, training, evaluation, deployment and monitoring. At MAS Holdings I led computer vision and IoT systems for manufacturing at scale, including care-label inspection that cut inspection time by 99.5% and raised operational efficiency by up to 300%. Your team's focus on taking models from research into reliable production systems is the work I have done throughout my career.

Three parts of your posting match my work directly. For the full model lifecycle, I owned the ML pipeline for the EU-funded MUSMET project from data acquisition and feature engineering through training, evaluation, deployment and monitoring. The published system runs at 14 ms on a Raspberry Pi 4, 74x faster than the DTW baseline, with F1 0.76 (Journal of the Audio Engineering Society, 2026). For evaluation and training decisions, I design benchmark suites with baselines and ablations to choose architectures under compute limits, and I released four open datasets on Zenodo. For MLOps, I have worked with Airflow, MLflow, Docker and Jenkins CI/CD, and built Airflow and Snowflake pipelines on AWS, and at MAS I introduced CI/CD and code review to speed up ML adoption across teams.

I should be direct about the gaps. I have not run distributed training with Horovod or Ray Train, used Spark or feature stores in production, or operated canary releases and A/B testing on a retail-scale serving platform. I have used SageMaker and Kubernetes at a lighter level, and my deployments have been edge and embedded rather than large cloud serving. I have built the same pieces with other tools, and I expect to pick up your stack quickly. Within the first months I would aim to contribute to training pipelines and deployment reliability.

I am a Canadian permanent resident, eligible to work without sponsorship, and available immediately. I am based in Toronto and open to relocating to Vancouver for the onsite requirement. I would welcome the opportunity to discuss how my experience in end-to-end ML systems can support the Enterprise Data & AI team at lululemon. My portfolio, publications and open datasets are available at nishal.xyz.

Sincerely,

Nishal Stanislaus Silva, PhD